

Concurrency

Processes vs. Threads

✓ Processes:

Like runners in separate lanes → independent but must manage shared resources.

✓ Threads:

Like a team working in one lane → share resources within a process.

Synchronization Primitives

✓ Mutex (Mutual Exclusion):

Like a single special water bottle → only one thread can use it at a time.

✓ Semaphore:

Like a limited number of stretchers → multiple threads can access, but only up to a limit.

Critical Sections, Race Conditions, and Deadlocks

✓ Race Condition:

Two threads try to access the same resource at the same time, causing unexpected behavior.

Example: Two runners reach for the same water bottle at the same time → water spills.

✓ Deadlock:

Two threads wait for each other indefinitely, preventing progress.

Example:

Team Member A holds one stretcher but needs another.

Team Member C holds the other stretcher but also needs one more.

Neither can proceed → Deadlock!

✚ Proper synchronization (mutexes, semaphores) prevents race conditions and deadlocks in multi-threaded systems.

Why Use Concurrency?

✓ Improves speed & efficiency → Don't waste time waiting, do other tasks.

✓ Leverages modern multi-core CPUs → Keeps cores busy for better performance.

✓ Enhances UI responsiveness → Prevents freezing by running background tasks.

Why Is Concurrency Hard?

✓ Even experienced engineers struggle to avoid race conditions, deadlocks, and crashes.

✓ Requires careful synchronization (mutexes, semaphores) to prevent conflicts.

Real-World Failure: Therac-25 Radiation Overdose

✓ Issue: A race condition in the software caused incorrect radiation doses.

✓ Cause:

The system failed to synchronize shared variables.

High-dose mode accidentally activated, leading to severe overdoses.

● Problem: Race Condition

✓ Threads t1 and t2 modify radiationLevel at the same time → No synchronization

✓ Since radiationLevel += 10 and radiationLevel -= 10 are not atomic operations, the final value may be unexpected or incorrect.

✓ Different runs might produce different results (unpredictable behavior).

✓ Solution: Mutex (std::mutex)

✓ std::lock_guard<std::mutex> lock(radiationMutex);

This locks the shared variable (radiationLevel) to prevent multiple threads from modifying it simultaneously.

✓ Ensures safe access to radiationLevel, preventing race conditions.

✓ Now, the final value of radiationLevel is always predictable and correct.

✚ Concurrency can be powerful but must be carefully managed to prevent catastrophic failures

Strategies for Handling Concurrency

① Strategy 1 - Be Slow (Prioritize Correctness over Speed)

✓ Example: Regression Testing for ML Models

1. Copy models to a staging directory
 - (a) Copy baseline models (10 min)
 - (b) Copy test models (10 min)
2. Compare outputs using the same inputs.
3. Optimization Tip: (a) & (b) can run in parallel if copying is a bottleneck.

Key Idea: Prioritize accuracy and avoid concurrency issues by not rushing operations.

② Strategy 2 - Isolate

✓ Keep shared resources separate to minimize conflicts.

✓ Instead of multiple threads sharing the same variable, give each thread its own copy.

③ Strategy 3 - Use Mutex (and Don't Forget to Unlock)

✓ Mutex (std::mutex) ensures only one thread modifies a resource at a time.

✓ Always unlock mutexes to avoid deadlocks.

```
std::mutex mtx;
void safeFunction() {
    std::lock_guard<std::mutex> lock(mtx); // Ensures mutex is
    unlocked automatically
    // Critical section
}
```

④ Strategy 4 - Use Tools

✓ Static Analysis Tools (e.g., ThreadSanitizer, Helgrind) detect concurrency bugs.

✓ Thread-safe libraries reduce manual handling of locks.

⑤ Strategy 5 - Put Comments

✓ Clearly document thread safety for each class/function.

✓ Examples of thread-safety comments: